COSC 6335 *"Data Mining"*

Problem Set 5 Fall 2024

Reviewing Data Mining Papers

Submitted by:

Mostafa Mushsharat / 2400966

Paper Title- Unsupervised Deep Subgraph Anomaly Detection (Zhao et al., 2022)

Authors- Zheng Zhang, Liang Zhao

Table of Contents

Summary and Contribution

The research paper is broadly under the category of Network Data Mining. To be more specific it falls under the domain of anomalous subgraph detection through graph mining. The topic of this paper addresses challenges and limitations in existing graph anomaly detection methods, especially in detecting subgraph-level anomalies.

The contribution of this paper is centered around three challenges:

i) Difficulty in obtaining sufficient training labels in an end-to-end manner without intensive supervision

ii) Difficulty in preserving both geodesic distance and topological similarity among nodes for representing subgraph anomaly

iii) Difficulty in quantifying the degree of being abnormal under arbitrary and unknown anomaly types.

The paper proposes a framework called AS-GAE- Anomalous Subgraph

Autoencoder. The framework involves developing a location-aware graph autoencoder to detect anomalous subgraphs. By using a reconstruction-based anomaly detection method, the model operates in an unsupervised manner, negating the need for labeled data and handling the challenge of an exponentially large subgraph search space. The location-aware autoencoder incorporates locational context by sampling a set of random anchor nodes and calculating the shortest distance from each node to these anchors. This allows for global positioning of nodes within the graph, enhancing the model's ability to maintain contextual relationships within subgraphs. By reconstructing the graph from low-dimensional embeddings, the model captures majority patterns while highlighting anomalies through higher reconstruction errors, effectively distinguishing anomalous regions within the graph.

The framework also addresses the challenge of quantifying the degree of anomaly. The model employs a super modular anomaly quantization function, ensuring that the anomaly score for a subgraph is at least as high as the sum of its individual partitions.

By integrating the supermodular neural network with the location-aware graph autoencoder, the framework operates in an end-to-end manner. This design allows the model to utilize reconstruction results for accurate subgraph-level anomaly detection.

The framework also introduces unsupervised and weakly supervised optimization strategies for tuning parameters in the location-aware graph autoencoder and the supermodular anomaly scoring function. The framework's unsupervised loss function prioritizes the assumption that most of the graph is normal, incorporating reconstruction error and anomaly penalization. When partial anomaly information is available, the model modifies the loss function to penalize subgraphs outside known anomaly regions.

Evaluation of the Contribution

Novelty

The paper's framework addresses constraints in labeled data availability, presenting both unsupervised and weakly supervised training methods, offering ways to detect anomalies in graphs where full anomaly labels may not be available or where only partial labeling can be assumed. While unsupervised methods for anomaly detection are not new, the paper's particular approach—focusing on subgraph-level detection using the location-aware graph autoencoder—is a modified application that appears original.

The use of a supermodular function to quantify subgraph anomalies builds on existing anomaly scoring but extends it to better capture the collective anomaly effect in subgraphs. This approach combines subgraph connectivity and individual anomaly scores, a novel combination aimed at accurately determining subgraph-level anomalies.

The authors introduce a location-aware graph autoencoder addressing limitations in traditional graph autoencoders that fail to capture subgraph context and node dependencies accurately. This adaptation could be seen as a novel approach, especially in unsupervised graph anomaly detection.

Technical Quality

The AS-GAE framework is tested on both synthetic and real-world datasets, which cover a range of subgraph anomaly types. This approach allows for evaluating AS-GAE's robustness in detecting anomalies across different network types, adding credibility to the results. Also by evaluating variants of AS-GAE with components removed (such as the location-aware encoder or supermodular function), the study demonstrates the contribution of each component, supporting the soundness of the framework design. However, the method uses randomly chosen anchor nodes to determine node locations within the graph structure. The random selection might lead to inconsistency if the anchors do not cover the network's unique structure, especially in very large graphs.

The authors indicate that they used both synthetic and real-world datasets, with most of the real-world datasets being widely used and accessible. They also mention a GitHub repository for the code, suggesting they intended to make the code replicable. However, the synthetic data generation process would require some setup and parameter adjustments that are not fully detailed in the paper.

The authors present a comprehensive evaluation of their approach's strengths, but their discussion of potential weaknesses is relatively limited. But the paper does implicitly acknowledge that selecting anchor nodes for capturing node location in the global graph structure is a critical step.

Potential Impact and Significance

This paper addresses a potentially impactful problem in the fields of machine learning, data mining, and network science, which suggests it is likely to be read and cited over time. Anomaly detection in graphs is crucial across various fields, such as cybersecurity (detecting fraud or malicious activities), healthcare (identifying disease outbreaks or novel patterns in patient data), social network analysis, finance, and telecommunications. Many of these applications extend beyond core machine learning or data mining and directly impact industries and sectors where networked data is needed. This cross-disciplinary relevance can lead to broader interest in and citations for the paper. The paper highlights that most existing work in graph anomaly detection has focused on node-level or edge-level anomalies, while subgraph-level anomalies have received less attention due to technical complexities. This focus can attract citations from researchers working on similar high-dimensional, structured data, especially as network-based applications continue to grow. Given that real-world anomaly detection often lacks labeled data, the model's capability in unsupervised and weakly supervised settings is a significant strength. The AS-GAE model's flexibility in detecting anomalies without labels, or with only minimal guidance, makes it potentially applicable to real-world settings, boosting its relevance and chances of being cited.

Clarity of Writing

The paper is organized in a clear, logical sequence. Each section builds on the previous one, making it easier for readers to follow the authors' argument. The use of examples, such as the illustrative disease outbreak network and fraudulent user social network, is helpful for understanding the applications of subgraph anomaly detection. The paper is generally well-written with an academic tone. I have not found any typos in the paper.

Strongest Points

Three strongest points

1. Novel framework- The paper introduces a novel AS-GAE framework that combines location-aware encoding with supermodular scoring to detect subgraph-level anomalies in an unsupervised manner.

2. Evaluation techniques- Experiments on synthetic and real-world datasets demonstrate AS-GAE's superior performance over existing methods, highlighting its robustness across various network types and anomaly scenarios.

3. Applicability to Real-World Problems- The method is relevant to real-world applications, such as fraud detection and disease outbreak monitoring, where subgraph anomaly detection is crucial and labeled data are often scarce.

Weakest Point

One weakest point

The random selection of anchor nodes in the location-aware encoder may lead to inconsistent embeddings, potentially impacting performance on larger or more complex networks.

Educational Value

The paper has strong educational value for graduate students interested in anomaly detection, graph neural networks, and unsupervised learning. It tackles a complex problem—subgraph anomaly detection—by introducing novel concepts like the location-aware autoencoder and supermodular anomaly scoring, which can deepen students' understanding of how to extend traditional machine learning methods to structured data. The paper's evaluation of AS-GAE against a range of existing methods provides valuable insights into benchmarking and the criteria needed to assess model performance in anomaly detection, making it a solid resource for students learning to design and evaluate machine learning experiments.

As an entry point to research in graph anomaly detection, the paper is informative but may require supplementary background reading. It outlines the unique challenges of subgraph anomaly detection, such as capturing complex subgraph structures without supervision, yet some concepts like supermodular functions might be challenging for students unfamiliar with advanced mathematical knowledge. Nevertheless, the authors' clear breakdown of each module and the role it plays in anomaly detection offers a practical framework for students to understand how such methods could be applied in real-world domains, like social networks or healthcare. This practical relevance may inspire students to explore how machine learning can be used to detect critical, real-world patterns that would be hard to identify manually.

The paper's emphasis on unsupervised and weakly supervised approaches, given the scarcity of labeled data in real-world scenarios, can also engage students interested in tackling data challenges. By showcasing AS-GAE's success across a variety of network types, it suggests that there is considerable room for improvement and adaptation in this field, especially as new graph neural network models and anomaly detection techniques continue to emerge. Overall, the paper has the potential to spark graduate students' interest in this research area, encouraging them to investigate further and potentially innovate upon the AS-GAE model or apply it to new datasets and applications.

From reading this paper, I learned about the different challenges and complexities involved in subgraph-level anomaly detection, especially within unsupervised settings where labeled data is not available. The authors provided insights into how graph-based anomalies are not simply nodes or edges but rather entire subgraphs with distinct structural or attributive properties, requiring sophisticated representation learning approaches. The AS-GAE model introduced new ideas, such as the location-aware graph autoencoder, which incorporates anchor-based location information to better capture the global context of nodes in the graph. Additionally, the concept of using a supermodular scoring function to assess the degree of anomaly at the subgraph level was insightful, as it highlighted the need to account for interdependencies within anomalies that appear in clusters rather than as isolated nodes or edges. Overall, the paper enhanced my understanding of both the theoretical aspects and practical applications of anomaly detection within graph structures.

Specific Comments and Questions

1.

Clarity on Anomaly Subgraph Definition: The definition of an anomaly subgraph is not clearly stated. While the paper provides examples of structural and attributive anomalies in Figure 1. *"…For example, the significant difference in subgraph topological structure is one type of anomaly subgraph. As shown in Figure 1(a), in a regular lattice material network, one might expect the inserted impurity as an outlier because the topological structure is different from other areas in the graph. Similarly, infectious disease researchers may have an interest in discovering a new unknown infectious disease at the early stage of a disease outbreak (e.g. COVID-19) from…", "…in red has a significantly different topology from other regions in the given network, which may be skeptical of a potential imperfection area. (b) An example of an attributive anomaly subgraph in a simulated disease outbreak network with three attributes of disease symptoms (fever, cough and broken bones). The potentially anomalous subgraph in the red dashed circle has an anomalous subset of attributes (cough and fever). The count number of these two attributes within the anomaly subgraph area are significantly higher than other nodes. when considered as a group compared to other areas. Merely aggregating node- or edge-level anomaly scores can not reflect…",* a formal definition would enhance clarity. For instance, what constitutes a "significant difference" in topological structure or attribute values? Defining specific thresholds or criteria would be helpful.

2.

Choice of Anchor Nodes: The location-aware graph autoencoder relies on sampling anchor nodes *"…To be specific, we first sample a set of C random anchor nodes VC = {v1, v2, . . . , vC} inthe given graph G. Then the shortest distance from each node to anchors, which is denoted as S ∈ R N×C , is calculated and treated as additional node labels.…".* However, the paper doesn't specify how these anchor nodes are selected. Clarifying the sampling strategy and its potential impact on performance is important. Does the choice of anchor nodes affect the detection of specific anomaly types?

3.

Threshold for Residual Graph Construction: The construction of the residual graph involves a threshold τ to filter out nodes based on their reconstruction error *"…then the residual graph R is constructed as: R = (VR, ER) ,VR = {vi|vi ∈ V & ri < τ}, (6) where ER is the corresponding set of edges to the set of nodes VR and τ is a threshold to filter out the nodes.…".* How is this threshold determined? Is it a fixed value or dynamically adjusted? Elaborating on the threshold selection process and its influence on the final detection results would be beneficial.

4.

Choice of Supermodular Function: The paper utilizes a deep graph supermodular neural network for anomaly quantification. While the provided example of *f(G) = |V| + |E|* is simple, how is the specific supermodular function chosen for the experiments? Does the choice of function impact the detection of various anomaly types? Discussing the selection process and potential alternatives would be insightful.

5.

Weakly Supervised Setting - Scale Determination: In the weakly supervised setting, the paper introduces a scale *Q* that encompasses the ground truth anomaly subgraph *"…The difference point is that there exists an anomaly subgraph in the given area of nodes Q and there are no anomalies outside this scale. Thus we modify the second term to a formulation, which first further…".* How is this scale determined in practice? Providing insights into the scale selection process and its influence on the algorithm's performance is essential.

6.

Computational Complexity: The paper analyzes the training time of AS-GAE *"…demonstrate that our proposed method effectively utilized the provided scale to further improve the anomaly detection power under our framework. D. Training time analysis The training time of AS-GAE under both unsupervised learning and weakly supervised learning settings are given in Table V with respect to a varying number of nodes and a fixed value of average node degree. The results indicate that the training time of AS-GAE increases almost linearly as the number of nodes in the graph increases. Furthermore,…".* However, discussing the computational complexity of the algorithm, particularly the supermodular function optimization, would be helpful. How does the complexity scale with graph size and feature dimensionality?

7.

Comparison with Other Anomaly Types: The paper focuses on structural and attributive anomalies. How would AS-GAE perform on other anomaly types, such as community-based anomalies or temporal anomalies? Discussing the generalizability of the framework to diverse anomaly types would strengthen the paper's impact.

Summary of Web Search Findings

Detecting unusual patterns in networks helps identify disease outbreaks, financial fraud, and suspicious social media activity. Networks contain both attribute data and connection patterns, which makes finding anomalies difficult. The OmniAnomaly model uses neural networks to find unusual patterns in time series data, which applies to network analysis (Su et al., 2019). Recent research shows that traditional methods struggle with irregular network structures (Ma et al., 2021).

New techniques have emerged to handle these challenges. The AM-GCN method analyzes both node characteristics and connections to find anomalies (Wang et al., 2020). Zhao et al. (2020) developed self-supervised systems to understand relationships in time series data. In financial fraud detection, Wang et al. (2019) created a neural network that learns from both labeled and unlabeled examples. Tang et al. (2018) applied these concepts to monitor structural health in buildings.

Current methods face several limitations. Graph Convolutional Networks struggle to combine node features with connection patterns (Ren et al., 2019). While many anomaly detection methods exist, few work well with network data (Gogoi et al., 2011). Naseer et al. (2018) highlighted similar issues when detecting network intrusions.

The research paper by Zhang and Zhao, directly addresses these challenges by proposing a novel deep learning framework called AS-GAE (Anomalous Subgraph Autoencoder). The paper aligns well with the gaps identified in my literature review and offers several innovative solutions:

1. The paper introduces a location-aware graph autoencoder that better preserves both geodesic distance and topological similarity among nodes - addressing a key limitation mentioned in the literature review about handling complex network structures.
2. The authors propose a supermodular graph scoring function to quantify anomalies without requiring prior knowledge of anomaly types - this directly addresses the challenge of interpretability and adaptability mentioned in the literature review.
3. The method achieves state-of-the-art results across multiple datasets, demonstrating significant improvements over both traditional and deep learning approaches - supporting the literature review's assertion that deep learning methods are becoming increasingly important in this field.

The mining of anomalous subgraphs in networks remains a rapidly evolving field that, despite significant progress, requires continued advancement in interpretability, adaptability, and empirical validation to develop more robust methods for diverse applications.

Broader Impact

This paper's novel approach to anomalous subgraph detection has significant potential to impact numerous critical real-world applications. In public health, it could revolutionize disease outbreak detection by identifying emerging clusters in health surveillance networks, enabling faster response to potential epidemics. In cybersecurity and finance, the method could help detect coordinated cyber-attacks, fraud rings, and money laundering schemes by identifying suspicious patterns in transaction and computer networks. The approach is particularly valuable in materials science, where it can detect structural anomalies and defects, and in social network analysis, where it can identify coordinated disinformation campaigns and malicious group activities.

The unsupervised and weakly supervised nature of AS-GAE makes it especially valuable for real-world applications where labeled data is scarce or expensive to obtain. This characteristic opens up possibilities in scientific discovery, where the method could help identify novel patterns in protein interaction networks, chemical compounds, or other complex systems without requiring prior knowledge of what constitutes an anomaly. The paper's methodological innovations, particularly the location-aware graph autoencoder and supermodular scoring function, provide a strong foundation for future research in graph-based anomaly detection while offering practical solutions that can be deployed in production environments.

Perhaps most importantly, this work bridges multiple research communities and creates opportunities for cross-disciplinary collaboration. It elegantly combines insights from graph theory, deep learning, and anomaly detection while connecting theoretical computer science concepts with practical applications. This interdisciplinary approach opens new research directions, including the integration with other deep learning architectures, extension to dynamic and temporal graphs, development of interpretable methods, and investigation of scalability improvements. By enabling collaboration between machine learning researchers and domain experts in fields like epidemiology, cybersecurity, and materials science, this work has the potential to accelerate innovation and create more effective solutions to real-world problems.

Justification for ICDM 2022 Best Paper Award

The paper's recognition with the ICDM 2022 Best Paper Award can be attributed to its exceptional combination of theoretical innovation, practical utility, and comprehensive validation. The paper addresses a fundamental challenge in network analysis - detecting anomalous subgraphs - with a novel approach that overcomes three critical limitations of existing methods: the need for labeled training data, the difficulty in preserving both local and global graph properties, and the challenge of quantifying anomalies without prior knowledge of their patterns. The authors' solution, combining a location-aware graph autoencoder with a supermodular scoring function, represents a significant advancement in unsupervised learning on graphs, particularly impressive given the complexity of subgraph-level analysis compared to simpler node or edge-level tasks.

What likely set this paper apart from other strong contributions is its balance of theoretical depth and practical applicability. The authors provide rigorous theoretical foundations with proofs for their supermodular properties while also demonstrating superior empirical performance across an extensive range of synthetic and real-world datasets. The paper's thorough ablation studies and sensitivity analyses reflect a deep understanding of the method's behavior and limitations. Furthermore, the work's potential impact across diverse applications - from disease outbreak detection to materials science combined with its ability to work with minimal supervision makes it particularly valuable for real-world deployment. This rare combination of theoretical elegance, practical utility, and comprehensive validation embodies the highest standards of data mining research, making it a worthy recipient of the Best Paper Award.

KDD-2012 Criteria Scores

1. **Novelty (1-7): 6/7**

   Proposes a novel framework (AS-GAE) for unsupervised/weakly supervised anomalous subgraph detection. Introduces innovative components like location-aware graph autoencoder and supermodular graph scoring. Clear differentiation from prior work in handling subgraph-level anomalies vs node/edge-level.

2. **Technical Quality (1-7): 6/7**

   Strong theoretical foundation with proofs for supermodular properties. Appropriate experiments on both synthetic and real-world datasets. Thorough ablation studies and parameter sensitivity analysis. Well-designed experimental setup with multiple baselines. Could have included more discussion on parameter selection guidelines.

3. **Potential Impact and Significance (1-7): 5/7**

   Addresses an important problem with applications in disease outbreak detection, fraud detection etc. Method is general and can handle different types of anomalies. Could have discussed more real-world deployment considerations.

4. **Clarity of Writing (1-7): 6/7**

   Well-structured with clear problem formulation. Good use of figures and examples to illustrate concepts. Technical content is presented logically. Minor grammatical issues but overall, very readable.

Additional scores:

5. **Educational Value (1-7): 6/7**

   Clear explanation of background concepts. Good literature review. Well-explained methodology with intuitive examples.

6. **Broader Impact (1-7): 5/7**

   Important applications in multiple domains. Could have discussed ethical considerations and limitations more.

7. **Overall Score (1-7): 6/7**

   Strong technical contribution with novel ideas. Well-validated through extensive experiments. Clear writing and presentation. Good balance of theory and practical results.

The paper presents a significant advancement in anomalous subgraph detection with strong technical merits and clear practical value. The main strengths are its novel methodology and comprehensive empirical validation. While there could be more discussion on practical deployment considerations and broader impacts, overall, it's a very strong contribution to the field.

Usefulness of Task 5

Task 5 is a valuable learning exercise that taught me how to systematically review research papers. By analyzing an award-winning paper using the KDD-2012 criteria, I learned to evaluate research across multiple dimensions - from technical content to real-world impact. The task helped me understand what makes research papers stand out and gave me practical experience in paper reviewing, a skill I'll need in my academic career. Having a structured framework for the review process was particularly helpful, as it showed me how to break down the evaluation of complex research into clear, manageable components. Through reading this paper, I learned about several important topics in data mining and machine learning: graph neural networks, autoencoders, anomaly detection in networks, and supermodular functions. The paper also taught me how unsupervised and weakly supervised learning can work when labeled data isn't available. I found it interesting how the authors combined these concepts to solve real problems like detecting disease outbreaks and finding defects in materials. The paper showed me how theoretical ideas in computer science can be applied to solve practical problems.

References:

1. Zhang, Z., & Zhao, L. (2022, November). Unsupervised deep subgraph anomaly detection. In *2022 IEEE International Conference on Data Mining (ICDM)* (pp. 753-762). IEEE.
2. Doncheva, N.., Assenov, Y.., Domingues, F. S.., & Albrecht, M.. (2012). Topological analysis and interactive visualization of biological networks and protein structures. Nature Protocols , 7 , 670-685 . http://doi.org/10.1038/nprot.2012.004
3. Wang, Xiao., Zhu, Meiqi., Bo, Deyu., Cui, Peng., Shi, C.., & Pei, J.. (2020). AM-GCN: Adaptive Multi-channel Graph Convolutional Networks. Proceedings of the 26th ACM SIGKDD International Conference on Knowledge Discovery & Data Mining . http://doi.org/10.1145/3394486.3403177
4. Naseer, Sheraz., Saleem, Y.., Khalid, S.., Bashir, Muhammad Khawar., Han, Jihun., Iqbal, M.., & Han, Kijun. (2018). Enhanced Network Anomaly Detection Based on Deep Neural Networks. IEEE Access , 6 , 48231-48246 . http://doi.org/10.1109/ACCESS.2018.2863036
5. Waheed, Abdul., Goyal, Muskan., Gupta, Deepak., Khanna, Ashish., Al-turjman, F.., & Pinheiro, P.. (2020). CovidGAN: Data Augmentation Using Auxiliary Classifier GAN for Improved Covid-19 Detection. Ieee Access , 8 , 91916 - 91923 . http://doi.org/10.1109/ACCESS.2020.2994762
6. Su, Ya., Zhao, Youjian., Niu, Chenhao., Liu, Rong., Sun, Wei., & Pei, Dan. (2019). Robust Anomaly Detection for Multivariate Time Series through Stochastic Recurrent Neural Network. Proceedings of the 25th ACM SIGKDD International Conference on Knowledge Discovery & Data Mining . http://doi.org/10.1145/3292500.3330672
7. Mitchell, N.., Nash, L.., Hexner, D.., Turner, A.., & Irvine, W.. (2016). Amorphous topological insulators constructed from random point sets. Nature Physics , 14 , 380 - 385 . http://doi.org/10.1038/s41567-017-0024-5
8. https://www.semanticscholar.org/paper/093bd7a7dd31a8eb4da371d012da5fd272bb96ca
9. Chen, Zekai., Chen, Dingshuo., Yuan, Zixuan., Cheng, Xiuzhen., & Zhang, Xiao. (2021). Learning Graph Structures With Transformer for Multivariate Time-Series Anomaly Detection in IoT. IEEE Internet of Things Journal , 9 , 9179-9189 . http://doi.org/10.1109/JIOT.2021.3100509
10. Peixoto, Tiago P.. (2013). Hierarchical block structures and high-resolution model selection in large networks. ArXiv , abs/1310.4377 . http://doi.org/10.1103/PhysRevX.4.011047
11. Wang, Daixin., Qi, Yuan., Lin, J.., Cui, Peng., Jia, Quanhui., Wang, Zhen., Fang, Yanming., Yu, Quan., Zhou, Jun., & Yang, Shuang. (2019). A Semi-Supervised Graph Attentive Network for Financial Fraud Detection. 2019 IEEE International Conference on Data Mining (ICDM) , 598-607 . http://doi.org/10.1109/ICDM.2019.00070
12. Tang, Zhiyi., Chen, Zhicheng., Bao, Y.., & Li, Hui. (2018). Convolutional neural network-based data anomaly detection method using multiple information for structural health monitoring. Structural Control and Health Monitoring , 26 . http://doi.org/10.1002/stc.2296
13. Lai, S.., Ruktanonchai, N.., Zhou, Liangcai., Prosper, O.., Luo, W.., Floyd, J.., Wesolowski, A.., Santillana, M.., Zhang, Chi., Du, Xiangjun., Yu, Hongjie., & Tatem, A.. (2020). Effect of non-pharmaceutical interventions to contain COVID-19 in China. Nature , 585 , 410 - 413 . http://doi.org/10.1038/s41586-020-2293-x

14. Chakraborty, Gouri., Park, In-Hyeok., Medishetty, R.., & Vittal, J.. (2021). Two-Dimensional Metal-Organic Framework Materials: Synthesis, Structures, Properties and Applications.. Chemical reviews . http://doi.org/10.1021/acs.chemrev.0c01049
15. Weng, H.., Liang, Yunye., Xu, Qiunan., Yu, Rui., Fang, Z.., Dai, X.., & Kawazoe, Y.. (2014). Topological node-line semimetal in three-dimensional graphene networks. Physical Review B , 92 , 045108 . http://doi.org/10.1103/PhysRevB.92.045108
16. Ma, Xiaoxiao., Wu, Jia., Xue, Shan., Yang, Jian., Zhou, Chuan., Sheng, Quan Z.., & Xiong, Hui. (2021). A Comprehensive Survey on Graph Anomaly Detection With Deep Learning. IEEE Transactions on Knowledge and Data Engineering , 35 , 12012-12038 . http://doi.org/10.1109/TKDE.2021.3118815
17. Chen, B.., Upadhyaya, N.., & Vitelli, V.. (2014). Nonlinear conduction via solitons in a topological mechanical insulator. Proceedings of the National Academy of Sciences , 111 , 13004 - 13009 . http://doi.org/10.1073/pnas.1405969111
18. Ren, Hansheng., Xu, Bixiong., Wang, Yujing., Yi, Chao., Huang, Congrui., Kou, Xiaoyu., Xing, Tony., Yang, Mao., Tong, Jie., & Zhang, Qi. (2019). Time-Series Anomaly Detection Service at Microsoft. Proceedings of the 25th ACM SIGKDD International Conference on Knowledge Discovery & Data Mining . http://doi.org/10.1145/3292500.3330680
19. Orús, R.. (2018). Tensor networks for complex quantum systems. Nature Reviews Physics , 1-13 . http://doi.org/10.1038/s42254-019-0086-7
20. Zhao, Hang., Wang, Yujing., Duan, Juanyong., Huang, Congrui., Cao, Defu., Tong, Yunhai., Xu, Bixiong., Bai, Jing., Tong, Jie., & Zhang, Qi. (2020). Multivariate Time-series Anomaly Detection via Graph Attention Network. 2020 IEEE International Conference on Data Mining (ICDM) , 841-850 . http://doi.org/10.1109/ICDM50108.2020.00093
21. Gogoi, Prasanta., Bhattacharyya, D.., Borah, B.., & Kalita, J.. (2011). A Survey of Outlier Detection Methods in Network Anomaly Identification. Comput. J. , 54 , 570-588 . http://doi.org/10.1093/comjnl/bxr026